

The Compound

Article6 — operating thesis, strategy, and accumulated signal.

Last updated: 2026-04-16

Core thesis

Paid verification work funds the capture of reviewer-grade decision traces. Those traces improve the product. The improved product makes the next verification more defensible, faster, and more valuable.

That is the loop. Everything else is downstream.

The loop (visual)

```
Paid verification job
↓
Revenue (fund next methodology encoding)
↓
Decision traces captured:
- which rules are hardest
- what rationale passes audit
- what evidence reviewers cite
- where methods are ambiguous
- where extraction is noisy
- where projects stall before finalize
↓
Product gets smarter:
- better default rationale suggestions
- smarter evidence type matching
- calibrated STAC eligibility heuristics
- auto-detection of weak rules
↓
Next verification is faster, more defensible, more valuable
↓
Higher price, more pilots, more data
↓
(compound)
```

Strategic framing

Article6 is **not**:

- a static SaaS tool
- a generic AI platform
- "AI for carbon"
- a consultancy deck with software attached

Article6 **is**:

- a verification factory that gets smarter from every paid review
 - the operating layer that lets VVBs run verification in a standardized, traceable way
-

Bottleneck we sell

Not "buy our platform."

Pay us per verification to remove the painful bottleneck in the review process.

The bottleneck: making project verification defensible, traceable, and faster without losing rigor.

Scarce asset

For frontier AI labs: compute (GPUs).

For Article6: **high-quality adjudication data from real reviewers on real rules.**

Includes:

- which evidence types actually satisfy which rules
- what "good support" looks like
- what reviewers reject
- where methods are ambiguous
- where extraction is noisy
- where projects stall before finalize

That is the dataset nobody gets from public docs alone.

Moat

Not:

- raw methodology PDFs
- generic evidence uploads
- another review UI

Moat:

- live verifier workflow access
- decision history
- rule-level evidence attachments
- finalize gating failures
- audit events
- cross-project patterns

Differentiated access to context feedback loops is where progress becomes legible and the business model improves.

Pricing ladder

- First 5 pilots: \$1,500-2,500 per full methodology verification
 - After 3 case studies showing 60-70% time savings: raise to \$5K
 - VVBs charge \$50-100K/engagement — we save them weeks
-

Wedge

Forestation/reforestation (ARR) verification for VVBs. Per verification. Defensible outputs.

Don't serve every stakeholder at once. Don't build 20 disconnected product lines. Own one strong wedge first.

The shared protocol (our "Iron Dome")

Standardize across reviews and projects:

- same rule state model
- same evidence attachment structure
- same audit events
- same finalize gates
- same export format
- same reviewer standards

Gives us:

- comparability
 - trust
 - lower reviewer training cost
 - portability across projects/methods
 - stronger buyer confidence
-

Customer co-building

Don't want distant "users." Want:

- 2-3 verifier design partners
- real projects
- real review sessions
- weekly workflow iteration
- product changes tied to their actual bottlenecks

Culture

- truth-seeking
- no fake verification
- no performative AI
- smallest real improvement
- reviewer trust over flashy UX
- product discipline over random features

Sources

- Anthropic founding investor interview (Anjney Midha, 2026-04-16)
- Three-agent roadmap synthesis (Hermes, GPT, Grok — 2026-04-15)
- VVB pricing research (~\$50-100K per engagement)